

Containerizing Applications with Dockerfiles

Pasquale Domenico Colaianni
DTU Biosustain, Research Software Engineering
2025-11-26

Challenges of sharing software with others

POV: Reviewer

1. Obtain your software
 - git clone or download a .zip
2. Scan for files and instructions
 - Is this a Python codebase or ... ?
 - Do I have all tools installed? Interpreters, compilers, Poetry, pipenv
 - Do I have the *correct* version of each?
3. Is my OS (Linux/Mac/Win) supported?
4. Will the software affect the integrity of my system?
 - A bug might cause the software to write in the wrong path

POV: You

1. Hope the reviewer can execute the software
 - Create an environment as close as possible to yours
2. Hope the reviewer is allowed (and inclined) to install any missing parts

Goals

Formally:

- Environment consistency
 - Ability to recreate the same software environment
- Dependency management
 - Better than relying on documentation only
- Isolation
 - The packaged app is not affected by the rest of the system
 - Example: permissions in Android and iOS

Informally:

- Plan so that the reviewer must execute as few actions as possible to reproduce your results
 - Reduce the amount of requirements

What is Docker?

Docker is a software for packaging and running your application in an isolated environment.

The result can be easily shared with others.

Docker comes with a command line tool.

Advantages

Docker is widely used in bioinformatics and software development because it addresses several challenges that are common in these fields:

- Reproducibility
- Portability
- Software dependencies management
- Scalability
- Collaboration
- Version control of environments
- Integration with other bioinformatics tools (e.g. Nextflow)

The only solution? [skip]

Virtual machines (images, snapshots of)

- The hardware must support emulation:
 - CPU and IO, including GPU, storage, networking, etc. (VT-x, AMD-V, VT-d, AMD-Vi)
- They emulate an entire operating system
 - and not just your application: more memory is required

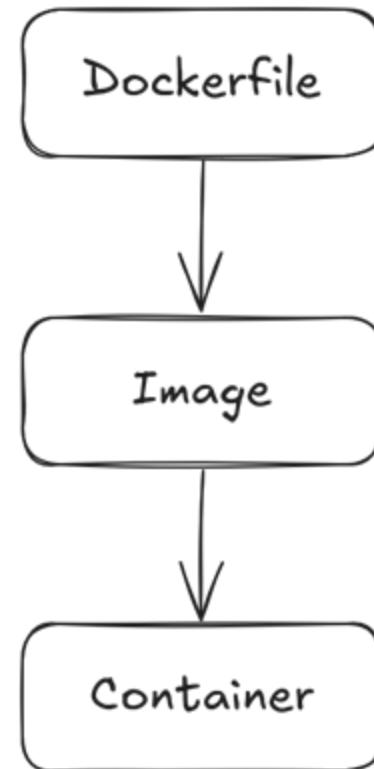
Docker containers

- They share the host's operating system
 - more memory available to your app
- Easier to package and share your work
- Faster to create, destroy, start, stop, etc.

Both solutions provide a consistent and isolated environment for your application.

Docker terminology

- Dockerfile
 - A text file that describes how to package and run your software
- Image
 - Created from a Dockerfile, with `docker build`
 - A package of your code and its dependencies
 - Is immutable: ideal for sharing with others
- Container
 - Created from an image, with `docker run`
 - A running instance of an image
 - Is mutable
 - Best if ephemeral, stateless



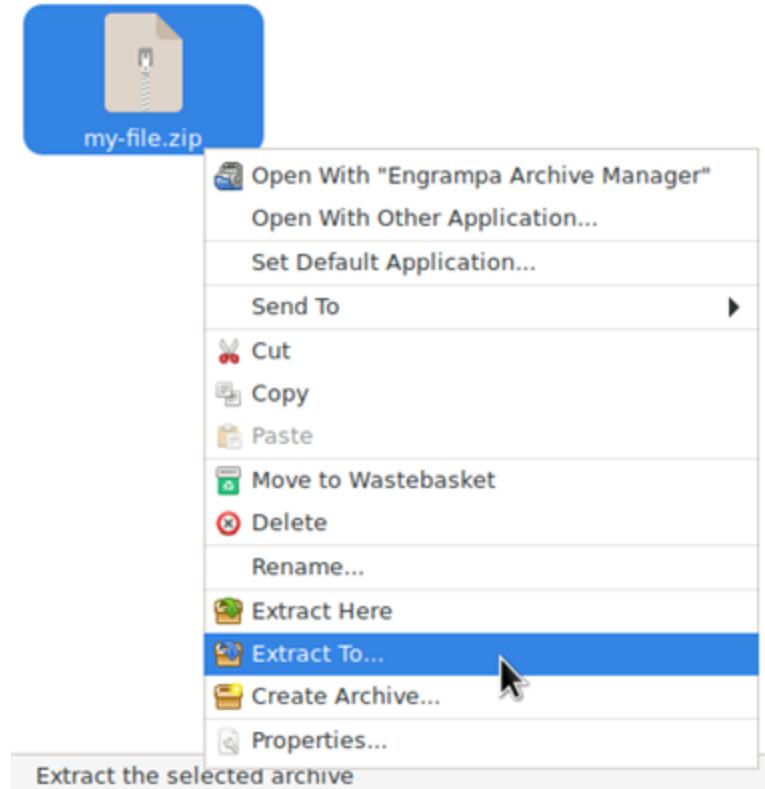
Glossary: <https://docs.docker.com/reference/glossary/>

Think of it like...

- Docker images as .zip files
- Docker containers as the extracted content

Point being: a .zip file can be extracted multiple times, to different destination folders. The content of those folders is the same.

People do not work directly on a .zip file. They work on the extracted content.



Bonus term: Registry

A storage and distribution system of Docker images. Docker Hub is the default registry (docker.io).

Registry terms:

- **repo name**: related images
- **tag**: label identifying a specific version of an image

Example:

```
docker pull docker.io/library/debian:12
```

```
docker pull debian:12
```

Dockerfile

It is a text file. It describes how to package and run your application.

```
FROM python:3.14
COPY main.py /app/main.py
CMD ["python", "/app/main.py"]
```

Find base images on Docker Hub:

- https://hub.docker.com/_/debian
- https://hub.docker.com/_/python
- https://hub.docker.com/_/ubuntu

No dependencies in this app.

How is the base image built? Where does it come from?

More instructions for Dockerfiles

- **ADD** - Add local or remote files and directories
- **ENTRYPOINT** - Specify default executable
- **ENV** - Set environment variables
- **RUN** - Execute build commands
- **WORKDIR** - Change working directory

All instructions: <https://docs.docker.com/reference/dockerfile/>

Images and containers - Commands

Given a Dockerfile, create an image:

```
docker build -t name[:tag] .
```

Given an image, create a container and start it:

```
docker run name[:tag]
```

Examples

https://github.com/biosustain/dsp_docker_training

Subfolder: *course_contents/Dockerfile_examples*

Exploring a container

```
docker pull debian:12
```

```
docker run --rm -it debian:12 bash
```

```
0 ~ docker pull debian:12
12: Pulling from library/debian
Digest: sha256:4abf773f2a570e6873259c4e3ba16
Status: Image is up to date for debian:12
docker.io/library/debian:12
```

```
0 ~ docker run --rm -it debian:12 bash
root@993b9e728972:/# whoami
root
root@993b9e728972:/# id
uid=0(root) gid=0(root) groups=0(root)
root@993b9e728972:/# ls /home
root@993b9e728972:/#
```

Enter `exit` to get out of the container. Option `--rm` removes the container.

Is docker pull needed?

```
~ docker run --help | grep -A 3 pull
  --pull string
          Pull image before running
          ("always", "missing", "never")
          (default "missing")
```

Running some commands...

... on the host:

```
+ hostname  
work-laptop  
  
+ ls -ahl /home  
total 12K  
drwxr-xr-x 3 root root 4.0K Nov 26 2023 .  
drwxr-xr-x 19 root root 4.0K Feb 8 22:46 ..  
drwx----- 42 user user 4.0K Feb 17 15:09 user  
  
+ ip -o link show  
lo  
eth0  
docker0  
br-123dfb4a82c1
```

... within the container:

```
+ hostname  
06dbdcf97766  
  
+ ls -ahl /home  
total 8.0K  
drwxr-xr-x 2 root root 4.0K Dec 31 10:25 .  
drwxr-xr-x 1 root root 4.0K Feb 17 14:14 ..  
  
+ ip -o link show  
ip: command not found
```

Try yourself with `docker run -it python:3.14 bash`

This underscores the isolation point made on slide 3.

Ask for help

```
docker build --help
```

```
docker run --help
```

```
docker ps --help
```

```
0 ~ docker run --help

Usage: docker run [OPTIONS] IMAGE [COMMAND] [ARG...]

Create and run a new container from an image

Aliases:
  docker container run, docker run

Options:
  --add-host list
  Add a custom host definition
```

Each help page provides a list of Aliases and Options.

Why Reproducibility Matters

- Scientific results must be verifiable by others
 - Inconsistent environments make that impossible
- Many bioinformatics tools behave differently depending on dependency versions
 - Which can lead to irreproducible or incorrect biological conclusions
- Regulatory, clinical, and high-stakes research settings require traceable and auditable computational pipelines
- Long-term studies need to re-run analyses years later
 - Without reproducible environments, workflows often break as software ecosystems change

Reproducibility

Tracking the versions of your dependencies is very important.

- `requirements.txt`
 - Basic reproducibility: it tracks direct dependencies versions
- lock file (`uv.lock`, `environment.lock`, `Pipfile.lock`, `poetry.lock`)
 - Good reproducibility: it tracks direct and indirects dependencies versions

Pick one:

- <https://docs.astral.sh/uv/> (strongly recommended)
- <https://docs.conda.io/projects/conda/en/latest/commands/index.html>
- <https://pipenv.pypa.io/en/latest/commands.html>
- <https://python-poetry.org/docs/cli/>

Reproducibility with conda [skip]

```
conda create --name app-with-conda python=3.12
conda activate app-with-conda
conda install conda-lock numpy
conda env export > environment.yml
conda-lock -f environment.yml -p linux-64
# To recreate the environment:
conda-lock install -n another-env -f conda-lock.yml
```

<https://www.anaconda.com/blog/8-levels-of-reproducibility>

Reproducibility with pipenv [skip]

```
pyenv local 3.13
python -m venv venv
source venv/bin/activate
pip install pipenv
pipenv install numpy
# It generated two files: Pipfile and
Pipfile.lock
```

- pyenv: <https://github.com/pyenv/pyenv>
 - Tool for installing and selecting versions of Python
- venv: <https://docs.python.org/3/library/venv.html>
 - Module from the standard library for creating virtual environments

```
FROM python:3.13
WORKDIR /app
COPY Pipfile.lock .
RUN pip install pipenv && pipenv sync
COPY app.py .
ENTRYPOINT ["pipenv", "run", "python", "app.py"]
```

Reproducibility in practice

- Track the lock file in your git repo
- Do not share your venv or secrets: use `.dockerignore` and `.gitignore`
- Keep your dependencies fresh
 - Once in a while (e.g. every 6 months) try upgrading your dependencies
 - If all works as expected, update the lock file in git
 - Avoid surprises due to breaking changes
 - pandas: <https://pandas.pydata.org/pandas-docs/stable/whatsnew/index.html>
 - numpy: <https://numpy.org/doc/stable/release.html>
- docker build --pull --no-cache
 - Use a fresh base image and no cache for your "production" images (images meant to be shared with others)

Exposing data to/from a container

- Via bind mounts
 - <https://docs.docker.com/engine/storage/bind-mounts/>
 - Alternative: volumes <https://docs.docker.com/engine/storage/volumes/>

```
FROM debian:12-slim
CMD ["cat", "/data/info.txt"]
```

```
0 cat-text-data tree
.
└── Dockerfile
    └── precious-data
        └── info.txt

2 directories, 2 files
0 cat-text-data docker build -q -t cat-text-data .
sha256:80b5adef402e8f178149dc77ed71b2dd74ec2b4d133cb48f65de7b9d43bd5f55
0 cat-text-data docker run cat-text-data
cat: /data/info.txt: No such file or directory
1 cat-text-data docker run -v $PWD/precious-data:/data:ro cat-text-data
I reside on the host's filesystem.
You can read me because the directory I reside
in is mounted to the container's filesystem.
0 cat-text-data
```

Network communication 1/3 [skip]

- Communication between containers
 - enabled by default for containers in the same bridge network
- Communication host <-> container
 - blocked by default; ports must be published

```
0 web-page tree
.
└── content
    └── index.html
└── Dockerfile

2 directories, 2 files
0 web-page docker build -q -t web-page .
sha256:82aef7d4ae23db3693563bde945a6681113066d63723759a
0 web-page docker run web-page
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
```

```
FROM nginx:stable
WORKDIR /usr/share/nginx/html
COPY content/index.html ./
```

Unable to connect

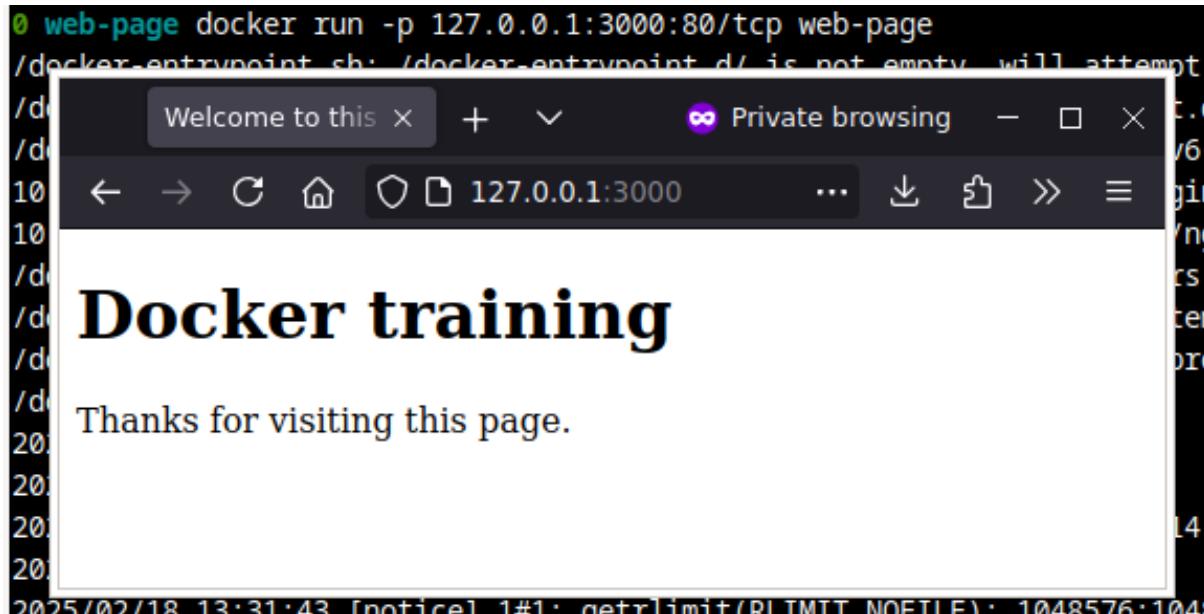
An error occurred during a connection to localhost.

- The site could be temporarily unavailable or too busy. Try again in a few moments.
- If you are unable to load any pages, check your computer's network connection.
- If your computer or network is protected by a firewall or proxy, make sure that Firefox is permitted to access the web.

Try Again

Network communication 2/3 [skip]

```
docker run -p host_ip:host_port:container_port/protocol
```



Challenge:
Can you achieve the same result, without copying `index.html` into the image?

<https://docs.docker.com/reference/cli/docker/container/run/#publish>

Network communication 3/3 [skip]

Challenge:

Can you achieve the same result, without copying index.html into the image?

Solution:

We could bind mount the folder containing the index.html file into the container.

```
docker run \
    -v $PWD/content:/usr/share/nginx/html:ro \
    -p 127.0.0.1:3000:80 \
    nginx:stable
```

Share Docker images

- Create an account on <https://hub.docker.com> and sign in
- Navigate to *Repositories*, create one
 - <https://hub.docker.com/repository/create>
- In your terminal, remember to execute `docker login`

Pushing images

You can push a new image to this repository using the CLI:

```
docker tag local-image:tagname new-repo:tagname  
docker push new-repo:tagname
```

Make sure to replace `tagname` with your desired image repository tag.

Multi-platform images

- Images can be built with support for multiple platforms
 - docker build --platform linux/amd64,linux/arm64 .
 - Provides support for the two most common architectures

<https://docs.docker.com/build/building/multi-platform/>

Authors might forget to build an image for your platform.

Solution: tell Docker which image to use!

docker run --platform linux/amd64 image-name:tag

```
0 ~ docker pull --help | grep platform
    --platform string      Set platform if server is multi-platform capable
0 ~ docker build --help | grep platform
    --platform stringArray  Set target platform for build
0 ~ docker run --help | grep platform
    --platform string      Set platform if server is multi-platform capable
```

Install instructions and best practices

- How to install: <https://docs.docker.com/engine/install/>
- Building images: <https://docs.docker.com/build/building/best-practices/>
 - Very important paragraphs:
 - Choose the right base image
 - Rebuild your images often
 - Exclude with .dockerignore
 - Create ephemeral containers
 - Don't install unnecessary packages
 - Sort multi-line arguments
 - Dockerfile instructions
 - Reference: <https://docs.docker.com/reference/dockerfile/>
- Code samples: <https://docs.docker.com/reference/samples/>
- Compose: <https://docs.docker.com/compose/how-tos/environment-variables/best-practices/>
- WSL: <https://docs.docker.com/desktop/features/wsl/best-practices/>

Some best practices covered today

- docker run --rm
- Read-only bind mounts
- Separate input and output bind mounts
- Changes within `docker exec` must be reflected in the Dockerfile
- Consider using docker compose

Further support

For Docker and software engineering: pasdom@dtu.dk

You're welcome to share feedback and ask for help.

Idea: software engineering hour, e.g. once a month, for topics such as git, programming languages, conventions, best practices, etc.